

Evolution of 3D Scene Representations

From Features to Neural Fields to Gaussians

Jugal Alan

M.Tech AI, IIT Ropar

2024AIM1004

Outline: The Evolution Journey

- ① **Distinctive Image Features from Scale-Invariant Keypoints (2004): SIFT**
- ② **Structure-from-Motion Revisited (2016): Modern 3D reconstruction**
- ③ **NeRF (2020): Neural radiance fields for photorealistic rendering**
- ④ **Instant-NGP (2022): Real-time neural graphics with hash encoding**
- ⑤ **3D Gaussian Splatting (2023): Explicit primitives for real-time rendering**

Key Question: How did we move from geometric pipelines to learned, real-time 3D representations?

Paper 1: SIFT - Scale-Invariant Feature Transform


David Lowe, IJCV 2004

Core Idea:

- Detect and describe local features
- Invariant to scale, rotation, illumination
- Enable wide-baseline matching


Impact:

- Backbone for SfM/MVS pipelines
- Enabled 3D reconstruction from images
- Foundation for modern feature matching



- Keypoints, → Matches

SIFT: Technical Pipeline



Key Components:

- **Difference of Gaussians (DoG):** Efficient scale-space extrema detection
- **Orientation histogram:** Achieves rotation invariance
- **Gradient-based descriptor:** 4×4 grid of 8-bin histograms = 128 dimensions
- **Matching:** Nearest-neighbor with ratio test to reject ambiguous matches

Why it matters: Reliable correspondences across viewpoints → enables 3D reconstruction

Paper 2: Structure-from-Motion Revisited

Schönberger & Frahm, CVPR 2016

Core Idea:


- Modern, robust incremental SfM
- Scales to thousands of images
- Produces camera poses + sparse 3D

Output:


- Camera extrinsics & intrinsics
- Sparse 3D point cloud
- Foundation for COLMAP system

Impact:

- De facto standard for pose estimation
- Provides input for NeRF training



SfM: Incremental Reconstruction Pipeline



Bundle Adjustment: Joint optimization minimizing reprojection error

$$\min_{\{\mathbf{R}_i, \mathbf{t}_i\}, \{\mathbf{x}_j\}} \sum_{i,j} \rho \left(\|\pi(\mathbf{R}_i \mathbf{x}_j + \mathbf{t}_i) - \mathbf{x}_{ij}\|^2 \right)$$

Robustness: RANSAC outlier rejection, next-best-view selection, redundant view mining

Transition: From Geometry to Neural Rendering

Traditional: SfM + MVS

- ✓ Accurate geometry
- ✓ Explicit structure
- ✗ Sparse reconstruction
- ✗ Limited photorealism
- ✗ Difficult novel views



Neural: NeRF

- ✓ Photorealistic rendering
- ✓ Continuous representation
- ✓ Novel view synthesis
- ✗ Slow training/rendering
- ✗ Requires posed images

Key insight: SfM provides the camera poses that NeRF needs. Now we can move from geometric reconstruction to learned radiance fields!

Paper 3: NeRF - Neural Radiance Fields

Mildenhall et al., ECCV 2020

Core Idea:


- Represent scene as continuous 5D function
- Map $(x, d) \rightarrow (\sigma, c)$

Revolutionary aspects:

- Photorealistic novel views
- No explicit geometry needed
- Continuous representation

Limitations:

- Slow training (hours-days)
- Slow rendering (seconds/frame)
- Per-scene optimization



NeRF: Technical Details

Input: Posed images (from SfM/COLMAP)

Representation: MLP $F_\theta : (\mathbf{x}, \mathbf{d}) \rightarrow (\sigma, \mathbf{c})$

- $\mathbf{x} = (x, y, z)$: 3D position
- $\mathbf{d} = (\theta, \phi)$: viewing direction
- σ : volume density (geometry)
- $\mathbf{c} = (r, g, b)$: emitted color (appearance)

Volume Rendering Equation:

$$\mathbf{C}(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \cdot \sigma(\mathbf{r}(t)) \cdot \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt$$

where transmittance $T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$

- Maps low-dimensional inputs to high-frequency features
- Enables MLP to represent fine details

NeRF: Training and Results

Training:

- Per-scene optimization (no generalization)
- Photometric loss: $\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \|\mathbf{C}(\mathbf{r}) - \mathbf{C}_{gt}(\mathbf{r})\|_2^2$
- Hierarchical sampling: coarse + fine networks
- Training time: hours to days per scene

Quality:

- State-of-the-art photorealistic novel view synthesis
- Captures complex view-dependent effects (reflections, specularities)
- Continuous representation allows arbitrary resolution

Bottleneck: Rendering requires querying MLP hundreds of times per ray → seconds per frame

Question: Can we keep the quality but make it fast?

Paper 4: Instant Neural Graphics Primitives

Müller et al., SIGGRAPH 2022

Core Idea:


- Replace large MLPs with hash encoding
- Store features in compact hash tables
- Small MLP for final prediction

Speed gains:

- Train in seconds to minutes
- Real-time rendering (interactive fps)
- 1000 \times faster than original NeRF

Impact:

- Enabled practical applications
- Foundation for many follow-ups



Multiresolution Hash Encoding

Instant-NGP: Hash Encoding Explained

Key Innovation: Multiresolution hash grids replace positional encoding + small MLP

How it works:

- ① Map 3D position \mathbf{x} to multiple resolution levels L_1, L_2, \dots, L_L
- ② At each level: hash \mathbf{x} to compact table, retrieve feature vector
- ③ Concatenate all features + $\mathbf{x} \rightarrow$ feed to tiny MLP (2 hidden layers, 64 neurons)
- ④ MLP outputs density σ and color \mathbf{c}

Hash function: $h(\mathbf{x}) = \left(\bigoplus_{i=1}^d x_i \pi_i \right) \bmod T$

- T : hash table size (e.g., 2^{19})
- π_i : large prime numbers

Why it's fast:

- Hash lookups are $O(1)$ and cache-friendly
- Small MLP \rightarrow fewer parameters, faster inference
- No expensive positional encoding computation

Instant-NGP: Performance Comparison

Method	Training Time	Rendering Speed	Quality (PSNR)
Original NeRF	1-2 days	30 sec/frame	31.0 dB
Instant-NGP	5-10 min	60+ fps	33.2 dB

Impact:

- Made NeRF practical for interactive applications
- Training time: days → minutes ($\sim 1000 \times$ speedup)
- Rendering: seconds → real-time ($\sim 100 \times$ speedup)
- Quality maintained or improved

From research curiosity to practical tool

Limitation: Still implicit representation → volume rendering → can we go faster with explicit primitives?

Paper 5: 3D Gaussian Splatting

Kerbl et al., SIGGRAPH 2023

Core Idea:


- Explicit 3D anisotropic Gaussians
- Differentiable splatting rasterization
- Adaptive density control

Breakthrough:

- Real-time rendering (≥ 30 fps at 1080p)
- Competitive quality with best NeRFs
- Editable, streamable representation

Advantages:

- Explicit primitives
- Memory efficient
- Enables editing/streaming



3D Gaussian Splatting: Representation

Each Gaussian is parameterized by:

- **Position:** $\mu \in \mathbb{R}^3$
- **Covariance:** $\Sigma \in \mathbb{R}^{3 \times 3}$ (anisotropic, full 3D ellipsoid)
- **Opacity:** $\alpha \in [0, 1]$
- **Color:** Spherical harmonics coefficients for view-dependent appearance

3D Gaussian function:


$$G(\mathbf{x}) = e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x}-\boldsymbol{\mu})}$$

Covariance representation: $\boldsymbol{\Sigma} = \mathbf{R} \mathbf{S} \mathbf{S}^T \mathbf{R}^T$

- **R:** rotation (quaternion)
- **S:** scaling (3D vector)
- Ensures positive semi-definite covariance

Initialization: From SfM point cloud (COLMAP output)

3D Gaussian Splatting: Rendering Pipeline



Key steps:

- ① **Projection:** Transform 3D covariance to 2D via Jacobian of perspective projection

$$\Sigma' = \mathbf{J} \mathbf{W} \Sigma \mathbf{W}^T \mathbf{J}^T$$

- ② **Tile-based rasterization:** Sort Gaussians by depth per screen tile (GPU-friendly)

- ③ **Alpha compositing:** Front-to-back blending

$$\mathbf{C} = \sum_{i=1}^N \mathbf{c}_i \alpha_i \prod_{j=1}^{i-1} (1 - \alpha_j)$$

3D Gaussian Splatting: Optimization

Loss function: $\mathcal{L} = (1 - \lambda)\mathcal{L}_1 + \lambda\mathcal{L}_{\text{D-SSIM}}$

Adaptive density control: Key to quality

- **Clone:** Large Gaussians in under-reconstructed areas
- **Split:** Large Gaussians with high gradients (need more detail)
- **Prune:** Gaussians with low opacity ($< \epsilon_\alpha$)
- Applied every N iterations during training

Training time: 30 minutes on RTX 3090 (vs days for NeRF)

Rendering speed: Real-time at 1080p (≈ 30 fps)

Method	Quality (PSNR)	Rendering Speed
Mip-NeRF 360	27.7 dB	0.3 fps
Instant-NGP	25.5 dB	6 fps
3D Gaussian Splatting	27.2 dB	138 fps

3D Gaussian Splatting: Trade-offs

Advantages:

- Real-time, high-quality rendering
- Explicit representation → editable, streamable
- Fast training (30 min vs days)
- Memory efficient for bounded scenes
- Enables applications: VR, games, robotics


Limitations:

- Gaussian count can grow (millions for complex scenes)
- Less naturally handles unbounded scenes than some NeRF variants
- Explicit primitives may have artifacts at boundaries
- Still requires SfM preprocessing (camera poses)

Extensions:

- 4D Gaussians for dynamic scenes (CVPR 2024)
- Gaussian Splatting SLAM for real-time mapping (CVPR 2024)

Evolution Timeline: The Full Picture




Key transitions:

- **Geometric era:** Features + optimization → accurate but sparse
- **Neural implicit era:** Continuous fields → photorealistic but slow
- **Explicit learned era:** Optimized primitives → real-time + quality

Quality vs Speed: The Trade-off Space

Quality (PSNR)



When to Use Each Method?

Method	Use When...	Avoid When...
SIFT + SfM	Need accurate geometry, camera calibration, sparse 3D	Need photorealistic rendering or dense reconstruction
NeRF	Prioritize quality over speed, offline rendering, research	Need real-time or interactive rates
Instant-NGP	Need balance of quality and speed, interactive preview	Need absolute best quality or real-time deployment
3DGS	Real-time rendering, VR/AR, games, robotics/SLAM	Extremely large unbounded scenes, limited memory

Practical recommendation:

Future Directions & Applications

Active research areas:

- **Generalization:** Move from per-scene to single model (e.g., pixelNeRF, generalizable NeRFs)
- **Dynamic scenes:** 4D Gaussians, D-NeRF, Neural Scene Flow
- **Large-scale:** City-scale reconstruction, satellite imagery
- **Sparse inputs:** Few-shot reconstruction, single image
- **Semantic understanding:** Combine with segmentation, object-level editing

Real-world applications:

- **Robotics/SLAM:** Gaussian Splatting SLAM, real-time mapping
- **VR/AR:** Immersive environments, photorealistic avatars
- **Film/VFX:** Virtual production, view synthesis
- **E-commerce:** Product visualization, virtual try-on
- **Cultural heritage:** Digital preservation, virtual museums

Summary: Key Takeaways

- ① **SIFT:** Local features enable wide-baseline matching
- ② **SfM:** Geometric reconstruction produces camera poses + sparse 3D
- ③ **NeRF:** Neural radiance fields achieve photorealistic novel views
- ④ **Instant-NGP:** Hash encoding brings NeRF to interactive speeds
- ⑤ **3DGs:** Explicit Gaussians enable real-time with quality

The evolution in one sentence:

From geometry to neural implicit to explicit learned primitives

Questions?

References

- ① Lowe, D. G. (2004). Distinctive image features from scale-invariant keypoints. *IJCV*.
- ② Schönberger, J. L., & Frahm, J. M. (2016). Structure-from-motion revisited. *CVPR*.
- ③ Mildenhall, B., et al. (2020). NeRF: Representing scenes as neural radiance fields for view synthesis. *ECCV*.
- ④ Müller, T., et al. (2022). Instant neural graphics primitives with a multiresolution hash encoding. *SIGGRAPH*.
- ⑤ Kerbl, B., et al. (2023). 3D Gaussian splatting for real-time radiance field rendering. *SIGGRAPH*.